

Quang Phu Nguyen

quangphunguyen1804@gmail.com | [linkedin.com/in/quangphunguyen](https://www.linkedin.com/in/quangphunguyen) | github.com/pdz1804

AI Engineer | LLM agents, RAG, agent governance | English: IELTS 6.5

EXPERIENCE

AI Engineer — FPT Software AI Center, Ho Chi Minh City Nov 2025 – Present

Agentic ERP Platform · *team of 15, stealth mode, ~1,000 users / ~100 DAU* Jun 2026 – Present

- Platform where agents act inside business systems; development centres on the **agent core** (harness, tools, prompts, memory, Skills). My scope: runtime, permissions and governance, multi-tenant admin, integrations.
- Built the runtime **default-deny permission layer** (a restricted agent cannot act outside its boundary), per-customer admin controls with audit trail, and usage/cost reporting; wrote the permission-model spec.
- Shipped the first release of user-authored **Skills** and team sharing of agents, skills and artifacts; built a third-party integration from scratch.
- Consolidated user and group management onto one source of truth; review peers on architecture boundaries, permission correctness and migration safety.
- Gave three internal AI4ALL sessions (AWS Strands agents, a portal for agent projects, agent-as-a-judge) and mentored an intern.

Healthcare Agentic Chatbot & Management Portal · *team of 10, ~100 users* Nov 2025 – May 2026

- Built the core **agent harness from day one** for four domains (recommendation, health & lifestyle, product education, tech support); **>90%** accuracy on a golden dataset from the tester, BA and PM.
- Cut typical per-query response latency **30s → 15s → 8s**: benchmarked Claude Haiku 3.5 / Sonnet 4 / Sonnet 4.5 on accuracy vs. answer quality, then replaced agent self-discovery with search-and-rank that narrows domain, tools and context.
- Solo-built the Management Portal beyond contract scope (usage/latency dashboards, knowledge approval, prompt management, automated tests, audit log); customer accepted POC-to-production **~2 weeks earlier** than on the previous contract.

Smart Product Recommendation System · *team of 10* Oct – Nov 2025

- Joined on arrival, before the team moved on to the agentic chatbot. Rule-based and LLM hybrid recommendation engine for a skin-health scanning device, with latency and quality monitoring, compliance checks and automated release gates.

AI Engineer Intern — FPT Software AI Center Jun 2025 – Oct 2025

- Built **M3ARAG**, a locally deployable multi-agent RAG over PDFs, HTML and Office files (LangGraph, Docling, ColPali), and a hybrid-search (BM25 + vector + semantic) blog system on Azure.
- Developed a Dual Attention Model that extracts technical keywords from company websites, plus a Company–Patent alignment pipeline (with The University of Tokyo).

Research Assistant — HCMUT Jun 2024 – Dec 2024

- Designed an unsupervised framework to build a Knowledge Graph with minimal domain-expert input, and evaluated LLMs for entity and intent extraction from raw Vietnamese documents.

EDUCATION

Ho Chi Minh City University of Technology (HCMUT) Ho Chi Minh City

B.Sc. Computer Science (Applied AI) — Graduated, Excellent, GPA 3.8/4.0 Sep 2022 – 2026

- Academic Incentive Scholarship (4/8 semesters), OISP Scholarship (3/8 semesters), Consolidation Prize at the Bach Khoa Innovation Contest.
- Honorable Mention, 2023 ICPC Vietnam Northern Provincial Programming Contest.

HONORS & CERTIFICATION

Top 10, AI Riser Vietnam 2026 (Google): Sách Của Em, 4th of 10 in audience vote Oct 2026

AI1 Young Talent, Level 3 (FPT Software) Sep 2026

“Best Team” award, FPT Americas (ST25) 2025

Anthropic: Model Context Protocol, Advanced Topics May 2026

TECHNICAL SKILLS

Agents & LLM: LangGraph, LangChain, AWS Strands, MCP, RAG, tool calling, agent evaluation

Backend & data: Python, FastAPI, PostgreSQL, Redis, Temporal, Qdrant, React/TypeScript

Cloud & ops: AWS (Bedrock AgentCore), GCP (GKE), Azure, Docker, Kubernetes, ArgoCD, OpenTelemetry